STA 325 Case Study Report

Sherry He, Megan Kolenski, Ian Le, Aaron Siminou

Introduction

This case study explores how three factors affect the clustering of inertial particles in turbulent flows such as clouds. The data come from computer simulations: in each run, our collaborators simulated particles moving in a turbulent flow, identified clusters of particles, and measured the volume of each cluster. The dataset has 89 simulation runs, each summarizing the distribution of particle cluster volumes at one (Re, Fr, St) setting. The three parameters take the following values:

- **Reynolds number** Re: 90, 224, 398. This measures how strong the turbulence is.
- **Froude number** Fr: 0.052, 0.3, ∞. This measures the strength of gravity. Smaller Fr means stronger gravity, and $Fr = \infty$ means no gravity.
- **Stokes number** St: 14 values between 0.05 and 3. This measures particle inertia. Particles with small St follow the flow closely, while particles with large St are heavier and respond more slowly.

Our key research objectives are:

1. Predicting the mean, standard deviation, skewness, and kurtosis of the distribution at a new (Re, Fr, St) observation
2. Explaining how each parameter affects the response

Our final model is as follows, where Re_{224}, Re_{398}, $Fr_{0.3}$, and Fr_{∞} are 0/1 indicators for each level (baseline $Re = 90$, $Fr = 0.052$).

$$\begin{aligned}\log \text{responses} = \beta_0 + \beta_1 \log St + \beta_2 (\log St)^2 + \beta_3 (\log St)^3 + \beta_4 Re_{224} + \beta_5 Re_{398} + \beta_6 Fr_{0.3} + \beta_7 Fr_{\infty} \\ + \beta_8 Re_{224} Fr_{0.3} + \beta_9 Re_{224} Fr_{\infty} + \beta_1 0 Re_{398} Fr_{\infty} + \varepsilon, \qquad \varepsilon \sim N(0, \sigma^2)\end{aligned}$$

A separate model is fit for each of the 4 response variables, so each one has its own coefficients. There is no $Re_{398} Fr_{0.3}$ term because the dataset has no observations at $Re = 398$, $Fr = 0.3$.

This model addresses both objectives. For prediction, it achieves a LOOCV MSE of 0.088 when averaged across the 4 response variables. For inference, the categorical terms describe differences among Re and Fr settings, while the polynomial terms allow the effect of St to vary across its range.

Table 1: Comparison of LOOCV MSE accross different models

Model	Parameters	Mean	Standard deviation	Skewness	Kurtosis
Baseline (additive)	6	0.0192	0.7338	0.3561	1.3252
Interaction Model 1	13	0.0104	0.3290	0.0715	0.2119
Interaction Model 2	9	0.0114	0.2565	0.0554	0.1656
Cubic Model	11	**0.0093**	**0.1377**	**0.0480**	**0.1567**
GAM	27	0.0124	0.3337	0.0639	0.1761

Methodology

Responses and predictors

Each observation in the dataset reports the first four raw moments of the response distribution, which we converted to central moments with `raw2central`. From these we computed the mean, standard deviation, skewness, and kurtosis of each distribution, which are the 4 desired response variables.

Modeling

We compared linear models with a generalized additive model (GAM), which fits a separate smooth curve in $\log St$ for each (Re, Fr) combination. The GAM had a higher LOOCV MSE than the cubic model for all 4 response variables (see the table below). We selected the cubic model because it had the lowest LOOCV MSE among the candidates while retaining a compact polynomial form.

We fitted 4 linear models for comparison. All models predict the log of each response. The baseline model has no interaction terms. Interaction Model 1 adds interactions of $\log St$ with both Re and Fr, as well as $Re \times Fr$. Interaction Model 2 adds only $Re \times Fr$. The cubic model adds $\log St^2$ and $\log St^3$ on top of Interaction Model 2. These candidates were guided by the EDA, which showed strong relationships with St, Re, and Fr.

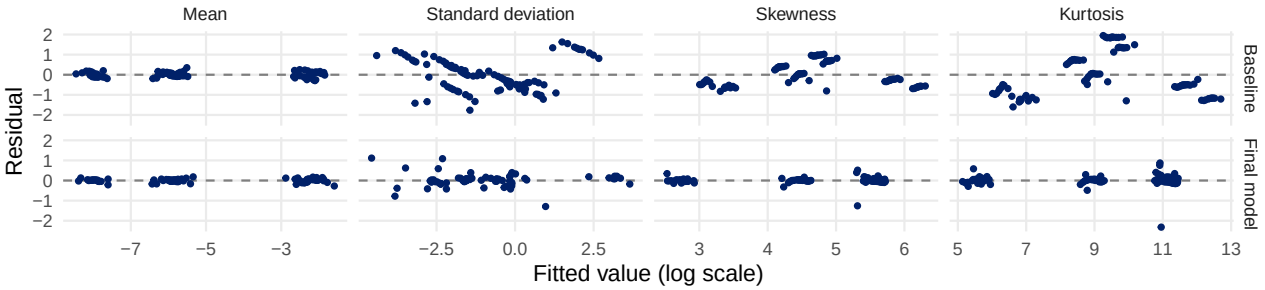
Averaged across the 4 response variables, the LOOCV MSE was 0.609 for the baseline model, 0.156 for Interaction Model 1, 0.122 for Interaction Model 2, and 0.088 for the cubic model. The cubic model had the lowest LOOCV MSE for all four response variables, so we selected it as our final model.

For the GAM, the parameter count is its effective degrees of freedom.

We log-transformed the response variables and St because they are right-skewed and span wide ranges. Re and Fr each have only 3 levels, so we treat them as categorical.

Diagnostics

The cubic model fits all 4 response variables well, with R^2 between 0.974 and 0.999. The plots below compare the residuals of the baseline model (top) and the cubic model (bottom).



The baseline residuals show clear patterns, especially for skewness and kurtosis, where whole groups of observations sit above or below 0. This means the effect of Re depends on Fr. The cubic model accounts for this interaction and curvature in St; its residuals are centered around 0, although some variation in spread remains, particularly for standard deviation.

The residual distributions still depart from normality, especially in their tails. The cubic terms allow the fitted relationship to bend across the observed St range, but observations at $St = 0.05$ remain worth monitoring because this is the boundary of the design. We kept these observations because there is no sign they are errors.

The dataset was collected from computer simulation runs, but observations with the same Re and Fr likely come from the same simulation, so they may not be fully independent. This could make our uncertainty estimates slightly too small.

Results

Predictions on the hold-out set

We predicted the four statistics at the 23 hold-out settings, with 95% prediction intervals computed on the log scale and then exponentiated, so point predictions estimate medians. All values are provided in `predictions.csv`. Every hold-out setting uses an (Re, Fr) combination and St value within the training range, so no prediction requires extrapolation. The two settings with $St = 0.05$ lie at the boundary of the observed range and should be interpreted with that limitation in mind.

The precision of the predictions varies across the four statistics. Across the hold-out settings, the upper bounds are about 1.4 times the lower bounds for the mean, 3.9 times for standard deviation, 2.2 times for skewness, and 4.1 times for kurtosis. The mean is predicted most precisely, while the standard deviation and kurtosis intervals are widest.

Effects of each parameter

Because the model is fit on the log scale, exponentiating a contrast gives the multiplicative change in a statistic. Since Re and Fr interact, the gravity contrast at a given Re combines the Fr coefficient with the corresponding interaction coefficient. The polynomial terms mean the effect of St varies across its range rather than remaining a constant multiplier.

Turbulence intensity (Re) controls cluster size. Under strong gravity, the cubic model estimates that increasing Re from 90 to 398 shrinks mean cluster volume by about 410 times and standard deviation by about 340 times. With no gravity, the reductions are about 250 times for the mean

and 9.6 times for standard deviation. Thus, Re has a large association with mean cluster volume. Under strong gravity, its estimated effects on skewness and kurtosis are small and not statistically distinguishable from zero in this model.

Strong gravity creates heavy tails, but only when turbulence is weak. At $Re = 90$, moving from strong gravity ($Fr = 0.052$) to no gravity ($Fr = \infty$) reduces skewness by a factor of about 16 and kurtosis by about 244. Strong gravity is therefore associated with a distribution having many small clusters and occasional very large ones. This effect weakens as Re increases: at $Re = 224$, removing gravity reduces kurtosis by about a factor of 9, and at $Re = 398$ it has little effect. One physical interpretation is that as turbulence strengthens, turbulent velocity fluctuations overwhelm gravitational settling, so the flow rather than gravity determines how particles cluster.

Moderate gravity behaves like no gravity. In the cubic model, the fitted effects for $Fr = 0.3$ and $Fr = \infty$ are similar across the four responses. Within the range studied, gravity is most strongly associated with clustering outcomes when it is strong.

Particle inertia (St) has a curved association with the responses. The cubic terms allow the relationship between St and each statistic to change in slope across the observed range. A single percentage change for doubling St would therefore be misleading; the fitted curves show the response-specific patterns. The model uses a common curve across (Re, Fr) settings, while allowing the levels of the curves to differ by those settings.

Conclusion

Each parameter shapes the cluster-size distribution in a different way. Turbulence intensity is strongly associated with cluster scale, while strong gravity is associated with heavier tails when turbulence is weak. The cubic model also indicates that the association with particle inertia changes across the observed St range. Because gravity's influence varies with turbulence strength, conclusions drawn at one turbulence level should not be assumed to hold at another. The model predicts mean cluster volume most precisely; uncertainty is wider for the other statistics. Its main limitations are the small design (three levels each of Re and Fr, with no runs at $Re = 398$, $Fr = 0.3$), remaining residual departures from normality and spread, and possible correlation within simulations. New simulations at intermediate Re and Fr values and at the missing combination would address these gaps most directly.

Appendix: Exploratory Data Analysis

Data Processing

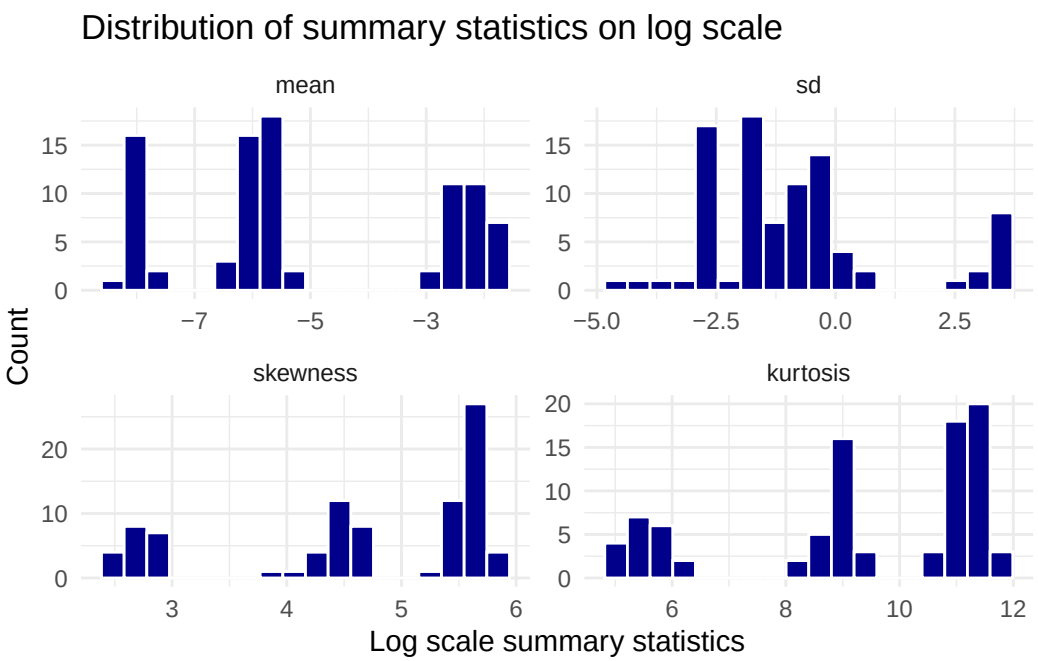
We see that the existing data contains Re and Fr values that are collected at distinct levels with low cardinality, instead of continuous values. It includes a valid infinity value for one of the predictors that is not a result of missing data. We compute standard deviation, skewness, and kurtosis from raw moments.

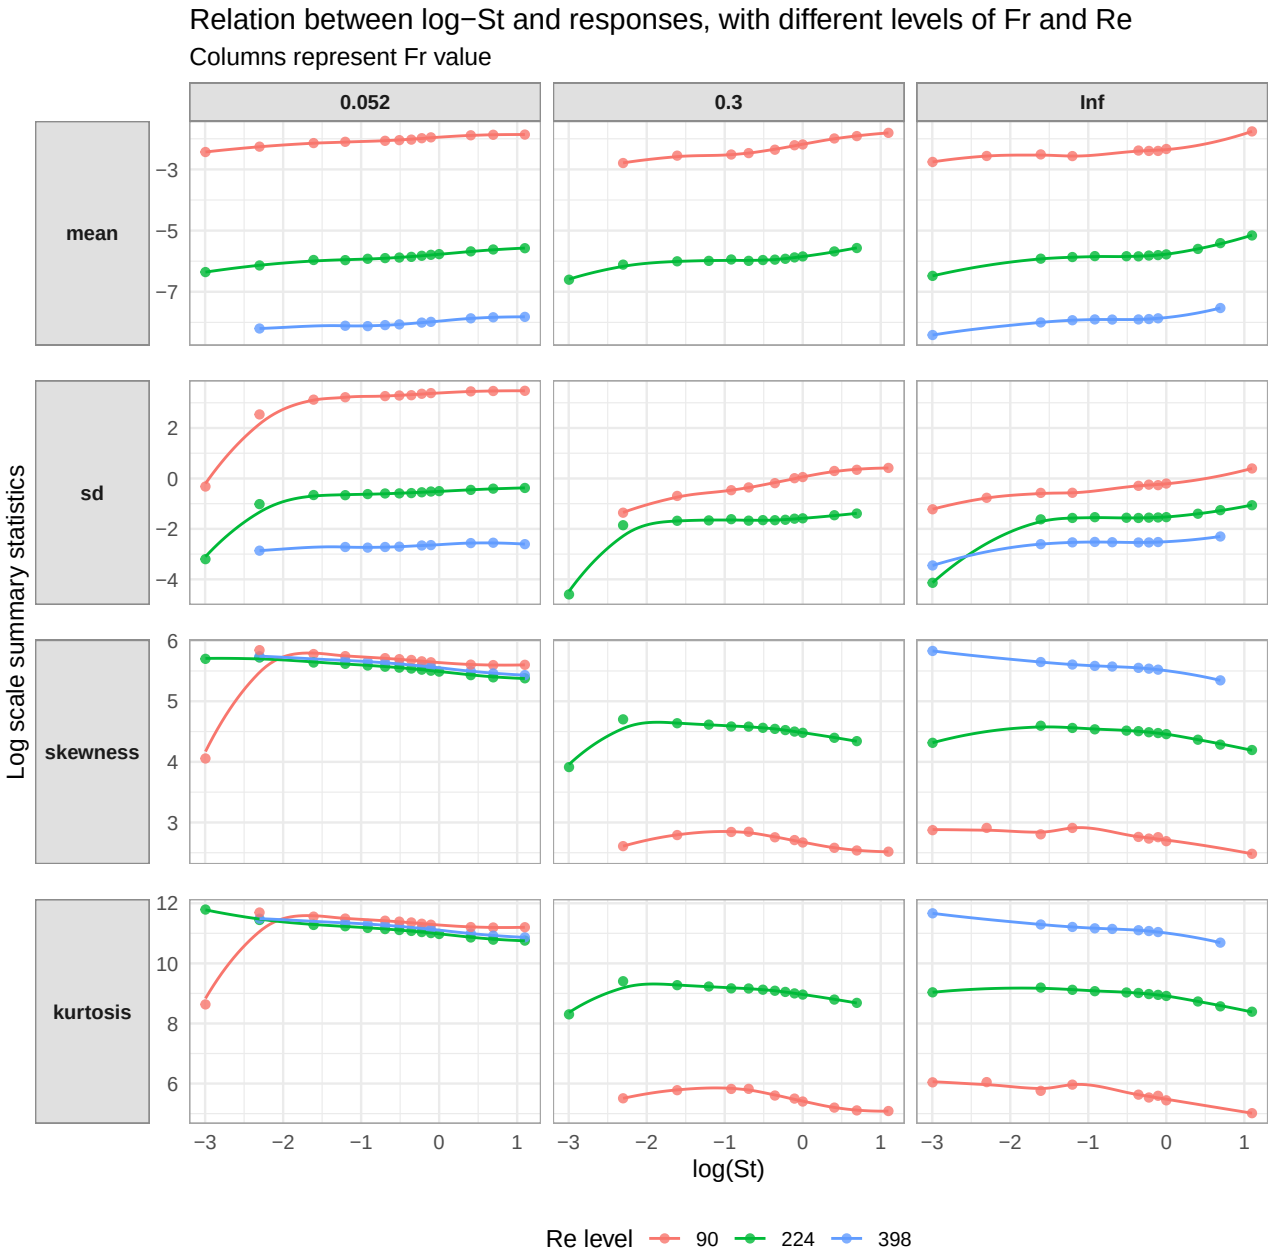
predictor	level	n
Re	224	39
Re	398	19
Re	90	31
Fr	0.052	36
Fr	0.3	23
Fr	Inf	30

response	min	q1	median	mean	q3	max	sd
mean	0.00	0.00	0.00	0.04	0.09	0.17	0.06
sd	0.01	0.16	0.28	3.65	0.73	32.32	8.95
skewness	11.97	72.55	110.12	162.81	269.54	344.91	113.10
kurtosis	150.53	5622.34	12158.72	39749.56	72732.45	132136.73	38100.12

Exploratory Data Analysis

Re	Fr_condition	n	St_min	St_max	St_levels
90	0.052	12	0.05	3	12
90	0.3	10	0.10	3	10
90	Inf	9	0.05	3	9
224	0.052	14	0.05	3	14
224	0.3	13	0.05	2	13
224	Inf	12	0.05	3	12
398	0.052	10	0.10	3	10
398	Inf	9	0.05	2	9


Distribution of summary statistics on log scale
mean
sd
skewness
kurtosis
Count
Log scale summary statistics


Relation between log−St and responses, with different levels of Fr and Re
Columns represent Fr value
0.052
0.3
Inf
mean
sd
skewness
kurtosis
Log scale summary statistics
log(St)
Re level
90
224
398